

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [19]: df=pd.read_csv(r"C:\Users\manasa\Downloads\fiat500_VehicleSelection_Dataset (3).csv")
df
```

Out[19]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
	0	1 lounge	51	882	25000	1	44.907242	8.611560	8900
	1	2 pop	51	1186	32500	1	45.666359	12.241890	8800
	2	3 sport	74	4658	142228	1	45.503300	11.417840	4200
	3	4 lounge	51	2739	160000	1	40.633171	17.634609	6000
	4	5 pop	73	3074	106880	1	41.903221	12.495650	5700


	1533	1534 sport	51	3712	115280	1	45.069679	7.704920	5200
	1534	1535 lounge	74	3835	112000	1	45.845692	8.666870	4600
	1535	1536 pop	51	2223	60457	1	45.481541	9.413480	7500
	1536	1537 lounge	51	2557	80750	1	45.000702	7.682270	5990
	1537	1538 pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

```
In [20]: df=df[['engine_power', 'age_in_days']]
df.columns=['eng', 'age']
```

```
In [21]: sns.lmplot(x="eng",y="age",data=df,order=2,ci=None)
```

```
Out[21]: <seaborn.axisgrid.FacetGrid at 0x16917519350>
```



```
In [22]: df.describe()
```

```
Out[22]:
```

	eng	age
count	1538.000000	1538.000000
mean	51.904421	1650.980494
std	3.988023	1289.522278
min	51.000000	366.000000
25%	51.000000	670.000000
50%	51.000000	1035.000000
75%	51.000000	2616.000000
max	77.000000	4658.000000

```
In [23]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    eng      1538 non-null    int64
1    age      1538 non-null    int64
dtypes: int64(2)
memory usage: 24.2 KB
```

```
In [24]: df.fillna(method='ffill',inplace=True)
```

C:\Users\manasa\AppData\Local\Temp\ipykernel_24456\4116506308.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df.fillna(method='ffill',inplace=True)

```
In [25]: x=np.array(df['eng']).reshape(-1,1)  
y=np.array(df['age']).reshape(-1,1)
```

```
In [26]: df.dropna(inplace=True)
```


C:\Users\manasa\AppData\Local\Temp\ipykernel_24456\1379821321.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df.dropna(inplace=True)

```
In [27]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)  
regr=LinearRegression()  
regr.fit(x_train,y_train)  
print(regr.score(x_test,y_test))
```


0.09511507783533979

```
In [28]: y_pred=regr.predict(x_test)  
plt.scatter(x_test,y_test,color='b')  
plt.plot(x_test,y_pred,color='k')  
plt.show()
```




```
In [29]: df500=df[:][:500]  
sns.lmplot(x="eng",y="age",data=df500,order=1,ci=None)
```

```
Out[29]: <seaborn.axisgrid.FacetGrid at 0x169175f4610>
```



```
In [35]: df500.fillna(method="ffill",inplace=True)
x=np.array(df500['eng']).reshape(-1,1)
y=np.array(df500['age']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

Regression: 0.08633689908802478



```
In [36]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.08633689908802478

```
In [ ]: #step-9 conclusion
#dataset we have taken is poor for linear model but with the smaller data works well with linear mode
```